Bitcoin Market Risk Analysis: Comparative Value-at-Risk Evidence from Hourly Data

Zhu Yecheng 3036061373
Jiang Sen 3036103860
Lin Fengtao 3035988603
Guo Yushan 3036103896

2026 April

Abstract

This study assesses the credibility of a Gaussian Value at Risk benchmark for Bitcoin when the asset is analyzed at an hourly frequency. The empirical design juxtaposes a rolling Gaussian benchmark against four more sophisticated model families: a parametric Student-t specification, a jump-diffusion Monte Carlo engine, a feature-conditioned LSTM, and a variational autoencoder employed for latent scenario synthesis. The data discredits the Gaussian baseline for both tail shape and backtest diagnostics, endorses heavy-tailed and state-dependent modeling, and demonstrates that the LSTM model is the sole specification to succeed in the conditional-coverage backtest at the 5% significance level.

Introduction

The main topic of the research is whether a simple Gaussian VaR framework can explain Bitcoin's downside risk over the course of an hour, or whether the data needs thicker tails, jump risk, and conditional machine-learning overlays. That question is important in academia because Bitcoin has big volatility clusters, sudden changes in events, and fast changes in regimes that can break the rules for textbook variance-scaling and thin-tailed parametric benchmarks (Jorion, 2007; McNeil, Frey, and Embrechts, 2015). The report assesses model credibility across three dimensions: distributional fit, out-of-sample tail calibration, and stress-period warning quality.

Data

The dataset has 46,279 cleaned hourly observations made from Binance BTC/USDT spot candles from January 1, 2021 to April 13, 2026. Data cleaning added 14 hours that were missing, got rid of 0 duplicate rows, and fixed 0 potential flash-crash artifacts. BTC/USDT is utilized as a liquid USD proxy, which makes it easier to sample every hour without interruption, but it should still be seen as a proxy benchmark rather than a pure fiat BTC/USD series.

Original rows	Clean rows	Missing hours	Duplicates removed	Flash-crash repairs	Start	End
46265	46279	14	0	0	2021-01-01 00:00:00 UTC	2026-04-13 06:00:00 UTC

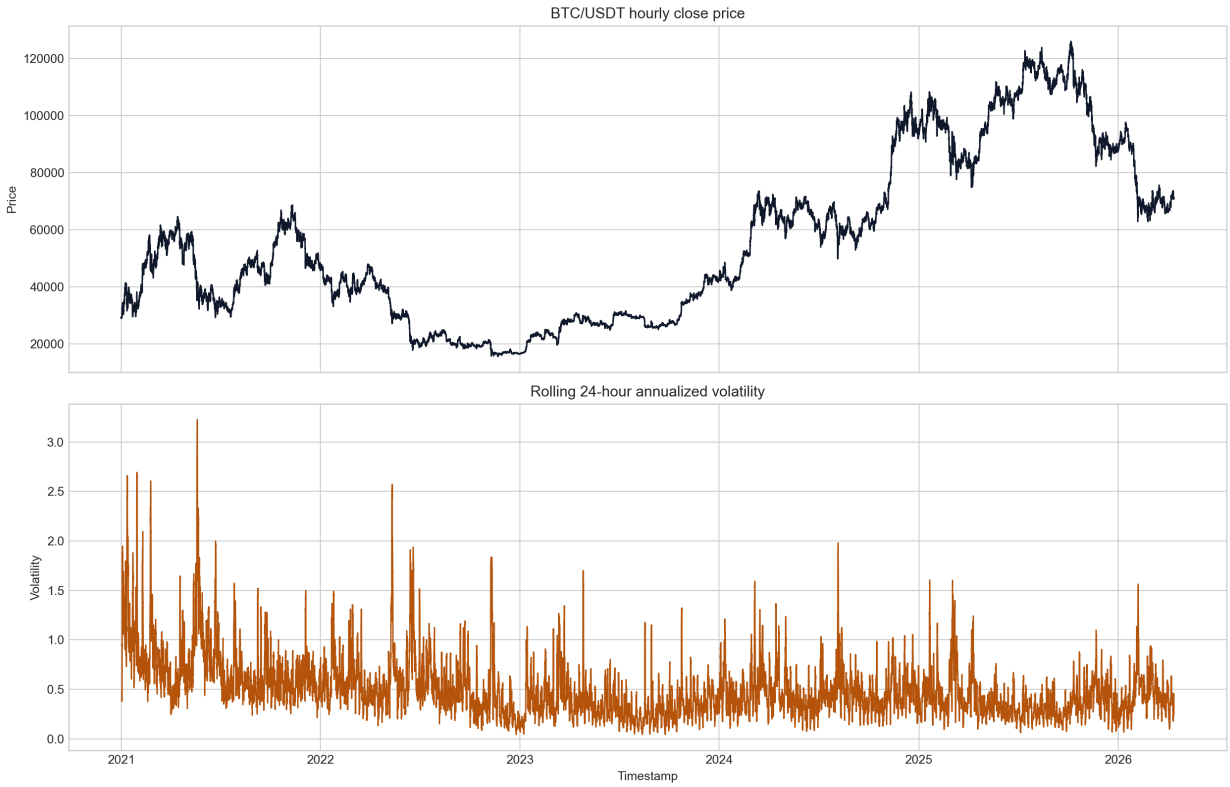

Figure 1: Price and volatility

Methodology

The benchmark layer contains two rolling parametric baselines: a Gaussian VaR benchmark and a Student-t VaR benchmark, both estimated on the same 90-day recent return window so that the distributional assumption is the main moving part (Jorion, 2007; McNeil, Frey, and Embrechts, 2015). The structural scenario layer adds jump-diffusion Monte Carlo with 10,000 simulation paths and a tuned jump threshold in a Merton-style discontinuous-return setting (Merton, 1976). The neural layer adds a one-step LSTM with a Gaussian output head and a VAE that learns latent return-window structure for synthetic scenario generation (Hochreiter and Schmidhuber, 1997; Kingma and Welling, 2014). Feature engineering supplies log returns, rolling 24-hour and 168-hour volatility, RSI(14), MACD with signal line, and a rolling volume z-score to the conditional models. Chronological splitting is enforced throughout. The final-year backtest holdout spans 2025-04-13 06:00:00 UTC to 2026-04-13 06:00:00 UTC with 8,761 hourly observations. Student-t and jump-diffusion hyperparameters are tuned on a 30-day pre-backtest slice sampled every 6 hours, which yields 120 validation checkpoints. For the neural models, the final 20% of the pre-backtest sequences or rolling windows are reserved for validation, giving 7,455 LSTM validation sequences and 7,498 VAE validation windows. The LSTM is trained with AdamW, learning rate 0.001, batch

size 256, and 8 epochs, with the final checkpoint chosen by minimum validation Gaussian negative log-likelihood. The VAE is trained with AdamW, learning rate 0.001, batch size 512, 20 epochs, and a reconstruction-plus-KL loss on rolling 24-hour windows. Validation scores in the tuning table therefore represent mean pinball loss for the parametric models, validation Gaussian NLL for the LSTM, and validation ELBO-style loss for the VAE.

Method Overview

Method	Model	Primary role	Main trade-off
A	Gaussian VaR	Thin-tailed rolling benchmark used to test whether a Gaussian assumption is adequate.	Transparent baseline, but structurally weak under heavy tails and volatility clustering.
B	Student-t VaR	Parametric heavy-tail benchmark using a rolling return window.	Transparent and fast, but only weakly conditional on current market state.
C	Jump-diffusion Monte Carlo	Scenario engine that combines continuous diffusion with discrete jump risk.	More realistic under stress, but sensitive to jump-calibration choices.
D	LSTM Gaussian head	Feature-conditioned one-step forecast of mean and volatility.	Adaptive to regime shifts, but harder to validate and explain.
E	VAE latent sampling	Generative model for synthetic return scenarios and latent tail structure.	Useful for scenario enrichment, but weaker as a standalone production VaR engine.

Tuned Model Settings

Model	Design	Best parameters	Selection metric	Validation score
Gaussian VaR	Rolling normal distribution benchmark	90-day rolling window	Baseline reference	
Student-t VaR	Parametric heavy-tail baseline	90-day rolling window; df estimated from sample kurtosis	Mean pinball loss	0.0003
Jump-diffusion Monte Carlo	Diffusion plus calibrated jump scenario engine	90-day rolling window; jump threshold 2.5 sigma	Mean pinball loss	0.0003
LSTM cond. VaR	Feature-conditioned recurrent forecaster with validation-calibrated tail	72-hour look-back; 2 layers; 48 hidden units; dropout 0.10; calibrated tail alpha 0.0125; calibrated tail z -2.662	Conditional-coverage-targeted fine-tune	0.9624
VAE latent VaR	Latent generative return model on rolling windows	24-hour window; latent dimension 8	Best validation ELBO	0.5177

Sample Split Summary

Block	Start	End	Observations	Design note
Full cleaned sample	2021-01-01 00:00:00 UTC	2026-04-13 06:00:00 UTC	46279	Hourly cleaned candles used for descriptive statistics, figures, and feature construction.
Pre-backtest return history	2021-01-01 01:00:00 UTC	2025-04-13 05:00:00 UTC	37517	History available for model estimation before the final one-year holdout.
Tuning validation slice	2025-03-14 06:00:00 UTC	2025-04-13 00:00:00 UTC	120	Chronological validation checkpoints every 6 hours for Student-t and jump-diffusion tuning.
Backtest holdout	2025-04-13 06:00:00 UTC	2026-04-13 06:00:00 UTC	8761	Final one-year out-of-sample window used for Kupiec coverage testing.
LSTM fit sequences	2021-01-11 00:00:00 UTC	2024-06-06 14:00:00 UTC	29823	Chronological 72-hour feature sequences used for recurrent model fitting.

Continued on next page

Block	Start	End	Observations	Design note
LSTM validation sequences	2024-06-06 15:00:00 UTC	2025-04-13 05:00:00 UTC	7455	Final 20% of pre-backtest sequence targets reserved for neural validation.
LSTM backtest sequences	2025-04-13 06:00:00 UTC	2026-04-13 06:00:00 UTC	8761	Out-of-sample sequence targets used for the one-step conditional backtest.
VAE train windows	2021-01-01 01:00:00 UTC	2024-06-04 19:00:00 UTC	29996	Rolling 24-hour windows used to estimate the latent generative model.
VAE validation windows	2024-06-03 21:00:00 UTC	2025-04-13 05:00:00 UTC	7498	Final 20% of pre-backtest rolling windows used for chronological latent-model validation.

The downside tail metrics used in the report follow the standard left-tail definitions:

$$CVaR_\alpha = \frac{1}{1-\alpha}\int_\alpha^1 VaR_u \, du$$

For unconditional coverage backtesting, the Kupiec proportion-of-failures statistic is:

$$LR_{POF} = -2\ln\left(\frac{(1-\alpha)^{T-x}\alpha^x}{(1-\hat{p})^{T-x}\hat{p}^x}\right), \qquad \hat{p} = \frac{x}{T}$$

where T is the number of forecasts, x is the number of VaR violations, and $\hat{p}$ is the observed violation rate. The report complements this with Christoffersen-style independence and conditional-coverage diagnostics (Kupiec, 1995; Christoffersen, 1998).

Statistical Diagnostics

The descriptive data definitively refutes Gaussianity. The Jarque-Bera statistic is 491079.15, with a p-value that is very close to zero. The skewness is -0.207, and the kurtosis is 18.955. The ADF statistic of -30.208 rejects a unit root in the return series, and the hourly returns range from -9.38% to 11.61%. When looked at conjunction, these tests show that heavy-tailed and conditional models are better than a thin-tailed homoskedastic baseline.

Obs.	Mean	Std.	Skew	Kurtosis	JB stat	JB p-val	ADF stat
46278	0	0.0064	-0.2066	18.9551	491079.1521	<0.0001	-30.2084

ADF p-val	ADF 1% crit	ADF 5% crit	ADF 10% crit	Min	Max	P1	P99
<0.0001	-3.4305	-2.8616	-2.5668	-0.0938	0.1161	-0.0193	0.0186

Model Results

Over the full sample, Bitcoin delivered a cumulative price change of 144.96%. In the current forecast snapshot, the one-hour VaR rankings differ materially by methodology, with the most conservative 99% estimate coming from Jump-diffusion MC at 1.63%. The Gaussian benchmark currently reports a one-hour VaR loss of 1.31%, versus 1.47% for the Student-t benchmark estimated on the same rolling history. The Student-t model remains the most interpretable heavy-tail baseline, the jump-diffusion engine provides the most stress-aware structural scenarios, and the LSTM contributes the most explicitly conditional one-step signal. Because the recurrent model is estimated as a one-step conditional forecaster, the 24-hour comparison table includes only models that directly produce multi-period distributions through aggregation or simulation.

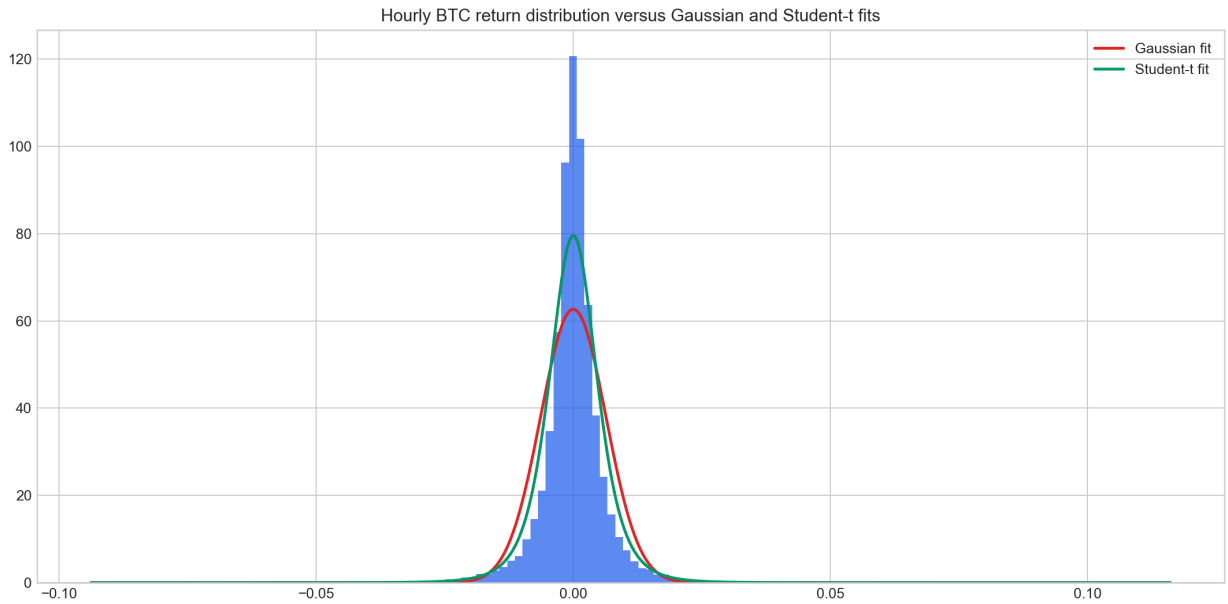

Figure 2: Return distribution

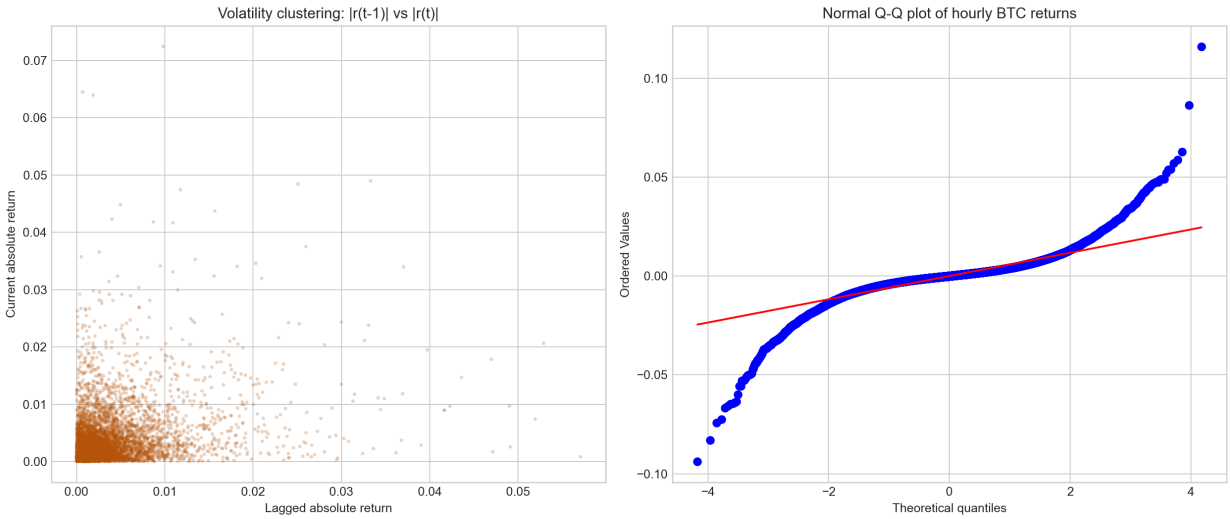

Figure 3: Distribution diagnostics

Model	Horizon (h)	VaR ret.	CVaR ret.	VaR loss	CVaR loss	Mean ret.	Vol.
Gaussian VaR	1	-0.0131	-0.015	0.0131	0.015	-0.0001	0.0056
Jump-diffusion MC	1	-0.0163	-0.0281	0.0163	0.0281	-0.0003	0.0066
LSTM	1	-0.0116	-0.0153	0.0116	0.0153	-0.0009	0.0043
Student-t param.	1	-0.0147	-0.0194	0.0147	0.0194	-0.0001	0.0043
VAE	1	-0.0117	-0.0144	0.0117	0.0144	0	0.0042
Gaussian VaR	24	-0.0665	-0.0757	0.0665	0.0757	-0.0029	0.0273
Jump-diffusion MC	24	-0.0829	-0.096	0.0829	0.096	-0.0046	0.0323
Student-t param.	24	-0.0672	-0.0779	0.0672	0.0779	-0.0029	0.0272
VAE	24	-0.0419	-0.0486	0.0419	0.0486	-0.0015	0.0174

Overall and Conditional VaR Summary

The overall VaR is measured from the full historical return distribution: the one-hour 99% empirical VaR loss is 1.93%, with a corresponding empirical CVaR loss of 2.79%. The conditional VaR requirement is met by the current one-hour forecasts; the most conservative current estimate is Jump-diffusion MC current forecast with a 99% VaR loss of 1.63% and CVaR loss of 2.81%. The LSTM row is also retained because it is explicitly conditioned on the latest feature state.

Scope	Basis	Horizon (h)	VaR loss	CVaR loss	Mean ret.	Interpretation
Overall empirical	Full-sample historical distribution	1	0.0193	0.0279	0.0000	Average unconditional one-hour tail risk over the full sample.
Conditional forecast	Jump-diffusion MC current forecast	1	0.0163	0.0281	-0.0003	Most conservative current one-hour forecast across model families.
Conditional forecast	LSTM current forecast	1	0.0116	0.0153	-0.0009	Feature-conditioned one-step forecast using the latest market state.

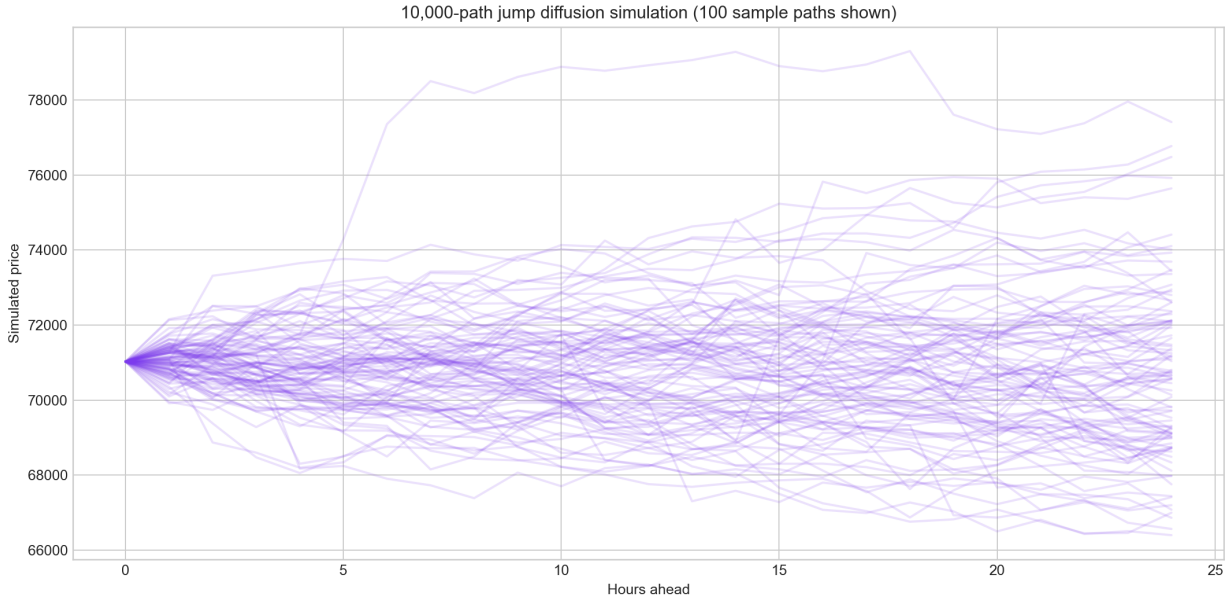

Figure 4: Monte Carlo paths

Sensitivity Analysis

Sensitivity analysis shows that the downside risk is likely to grow when confidence rises from 95% to 99%. The biggest one-hour VaR decrease is still 1.63% at 99%. The Gaussian row in the confidence table gives the thin-tailed benchmark, the Student-t parameter sweep illustrates that lesser degrees of freedom make tail losses worse in a nonlinear way, and the LSTM look-back sweep shows that you have to choose between responsiveness and stability.

Confidence-Level Comparison

Conf. (%)	Model	VaR loss	CVaR loss	Vol.
95	Jump-diffusion MC	0.0097	0.0143	0.0067
95	Gaussian VaR	0.0093	0.0116	0.0056
95	Student-t param.	0.0088	0.0126	0.0043
95	LSTM	0.008	0.0098	0.0043
95	VAE	0.0074	0.0101	0.0042
97.5	Jump-diffusion MC	0.0119	0.0179	0.0065
97.5	Student-t param.	0.0112	0.0153	0.0043
97.5	Gaussian VaR	0.0111	0.0132	0.0056
97.5	LSTM	0.0093	0.011	0.0043
97.5	VAE	0.0093	0.012	0.0042
99	Jump-diffusion MC	0.016	0.0284	0.0067
99	Student-t param.	0.0147	0.0194	0.0043
99	Gaussian VaR	0.0131	0.015	0.0056
99	VAE	0.0117	0.0144	0.0042
99	LSTM	0.0109	0.0124	0.0043

Student-t Degrees-of-Freedom Impact

DoF	VaR ret.	VaR loss	Tail ratio / sigma	Estimated fit
4.5	-0.0148	0.0148	2.6506	0
5	-0.0147	0.0147	2.6281	0
6	-0.0144	0.0144	2.5876	0
8	-0.0141	0.0141	2.5301	0
12	-0.0138	0.0138	2.4691	0
20	-0.0135	0.0135	2.4199	0

LSTM Look-back Window Impact

Look-back (h)	Best val. loss	Sigma MAE	Current VaR loss	Current sigma	Val fail rate	Val POF p	Val IND p	Val CC p	Tail z
24	-4.6994	0.0038	0.014	0.0054	0.01	0.9675	0.0473	0.1397	-2.2808
48	-4.7665	0.0033	0.0123	0.0049	0.0101	0.9629	0.7852	0.9625	-2.5731
72	-4.7694	0.003	0.0124	0.0043	0.0101	0.9583	0.7857	0.9624	-2.662

Stress Testing

The stress section is designed as a two-window case-study exercise rather than an exhaustive event screen; one downside exchange-failure regime and one policy/news-driven whipsaw regime were

chosen to contrast crash risk and two-sided event volatility. Two pre-specified event windows are defined in code, and the event timestamp inside each window is selected algorithmically by either the worst one-hour loss or the largest absolute one-hour move; warnings are then evaluated over the preceding 24-hour horizon. Within that design, the FTX collapse window peaks at 2022-11-08 17:00:00 UTC with a realized one-hour loss of 5.18%, and the ETF whipsaw peaks at 2024-01-11 15:00:00 UTC with a realized absolute move of 3.75%. In both cases, the jump-diffusion engine issues the strongest warning profile and lands closest to the realized loss. At the same time, every one-hour VaR model understates the realized event loss by a wide margin, so the stress section should be interpreted as a warning-comparison exercise rather than a claim of exact crisis prediction.

Event	Event time	Realized return	Realized loss	t-VaR loss	MC VaR loss	Closest model
FTX Collapse	2022-11-08 17:00:00 UTC	-0.0518	0.0518	0.0145	0.0163	Jump-diffusion MC
ETF whipsaw	2024-01-11 15:00:00 UTC	-0.0375	0.0375	0.0123	0.0129	Jump-diffusion MC

Event	t peak warning	MC peak warning	t peak lead (h)	MC peak lead (h)	Strongest warning
FTX Collapse	0.0144	0.0171	5	23	Jump-diffusion MC
ETF whipsaw	0.0122	0.014	2	23	Jump-diffusion MC

Figure 5: Stress warning profiles

Time-Scale Scaling

The square-root-of-time rule doesn't always work for Bitcoin. The scaled one-day 99% VaR is 9.43%, while the empirical one-day VaR is 8.74%. The scaled ten-day VaR is 29.83%, while the

empirical ten-day VaR is 27.66%. Weak linear autocorrelation in raw returns appears alongside persistent autocorrelation in absolute returns, indicating that volatility clustering, rather than mean predictability, is the primary factor causing naive variance scaling to misrepresent long-horizon risk.

Horizon	Scaled VaR	Empirical VaR	Scaling bias %
1 day	0.0943	0.0874	7.9691
10 days	0.2983	0.2766	7.8213

r(1)	r(24)	|r|(1)	|r|(24)	LB p-value r(24)	LB p-value |r|(24)
-0.0047	-0.0262	0.2847	0.1904	<0.0001	<0.0001

Practical Implications

For a hypothetical $1 million BTC spot position, the largest 24-hour VaR capital estimate in the current snapshot is $82,946.43. These are market-risk numbers rather than liquidity-adjusted exit costs. For committee reporting, the jump-diffusion estimates are the more conservative baseline, while the LSTM is more appropriate as an intraday surveillance overlay.

Notional USD	Model	Horizon (h)	VaR capital	CVaR capital
1000000	Jump-diffusion MC	1	16276.37	28112.64
1000000	Student-t param.	1	14655.86	19350.76
1000000	Gaussian VaR	1	13094.56	14984.39
1000000	VAE	1	11732.93	14385.96
1000000	LSTM	1	11576.58	15282.24
1000000	Jump-diffusion MC	24	82946.43	96006.01
1000000	Student-t param.	24	67183.15	77946.75
1000000	Gaussian VaR	24	66455.78	75714.02
1000000	VAE	24	41882.46	48609.76

Backtesting

The last-year backtest uses both Kupiec's unconditional coverage test and Christoffersen's independence and conditional-coverage diagnostics to look at 8,761 one-hour forecasts against a 1% tail target (Kupiec, 1995; Christoffersen, 1998). The LSTM model has a conditional-coverage p-value of 0.078, while the Gaussian benchmark has a Kupiec p-value of 0.000 and a conditional-coverage p-value of 0.000. The LSTM model is the only one that passes the conditional-coverage test at the 5% significance level. That trend shows that conditional adaptation is a key feature in tail calibration on this holdout window. However, it is still evidence regarding this sample route and not proof that one specification works better in every environment.

Model	Obs.	Viol.	Exp. fail rate	Obs. fail rate	POF p-value	IND p-value	CC p-value
Gaussian VaR	8761	184	0.01	0.021	<0.0001	<0.0001	<0.0001
Student-t VaR	8761	129	0.01	0.0147	<0.0001	0.0001	<0.0001
Jump-diffusion MC	8761	110	0.01	0.0126	0.0208	0.0005	0.0002
LSTM cond. VaR	8761	72	0.01	0.0082	0.0837	0.1463	0.078

Continued on next page

Model	Obs.	Viol.	Exp. fail rate	Obs. fail rate	POF p-value	IND p-value	CC p-value
VAE latent VaR	8761	133	0.01	0.0152	<0.0001	0.0051	<0.0001

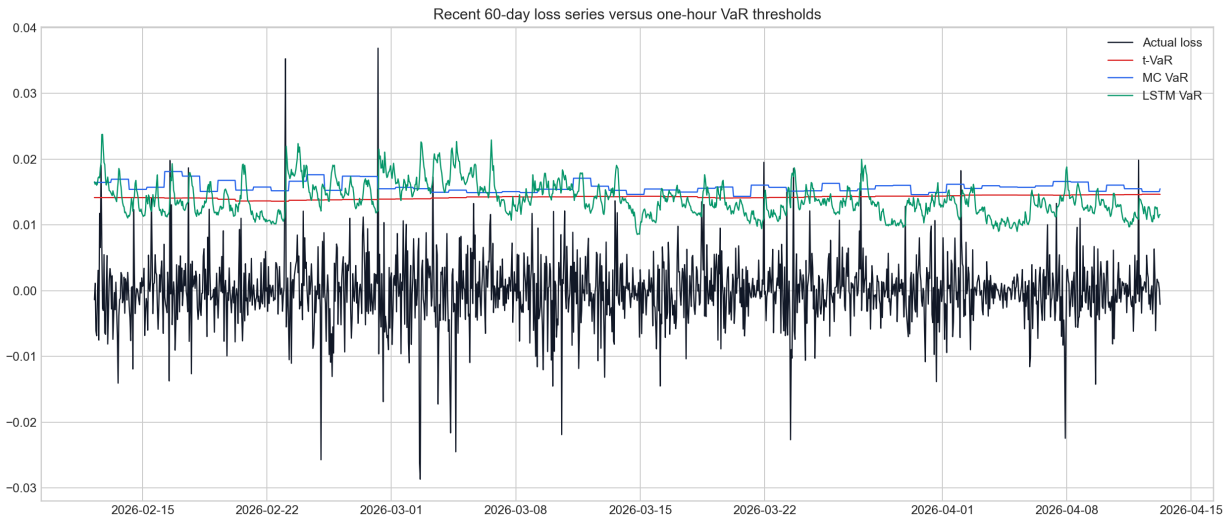

Figure 6: VaR backtest

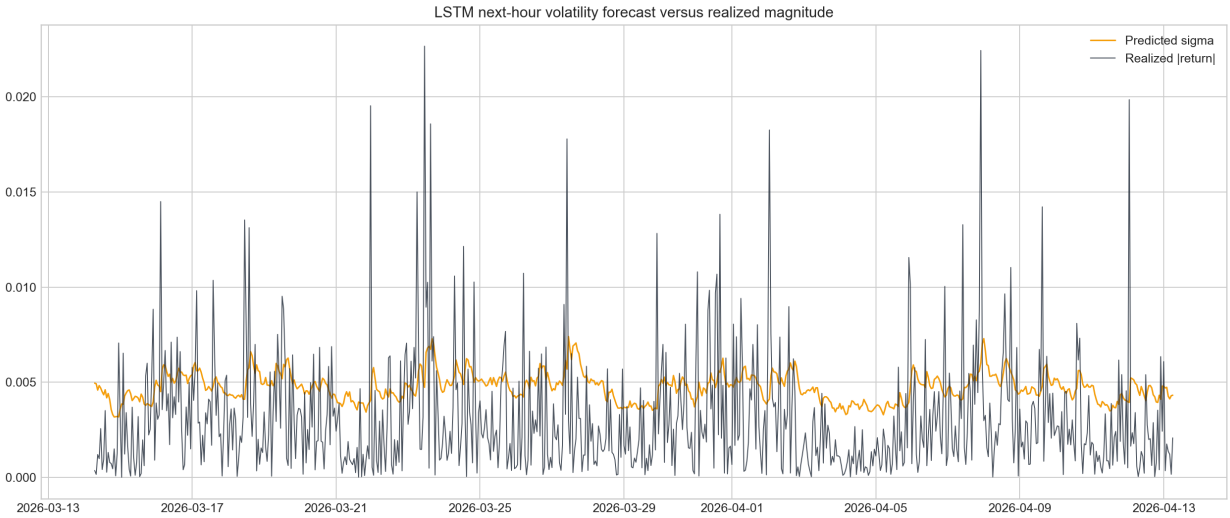

Figure 7: LSTM volatility forecast

Limitations

There are a few clear limits to the project. The benchmark is a proxy BTC/USDT series instead of a fiat BTC/USD market. The report uses hourly candles instead of intrahour depth, and the backtesting part only looks at one realized holdout year, even after adding conditional-coverage diagnostics. The neural models make things more flexible, but they are still harder to understand

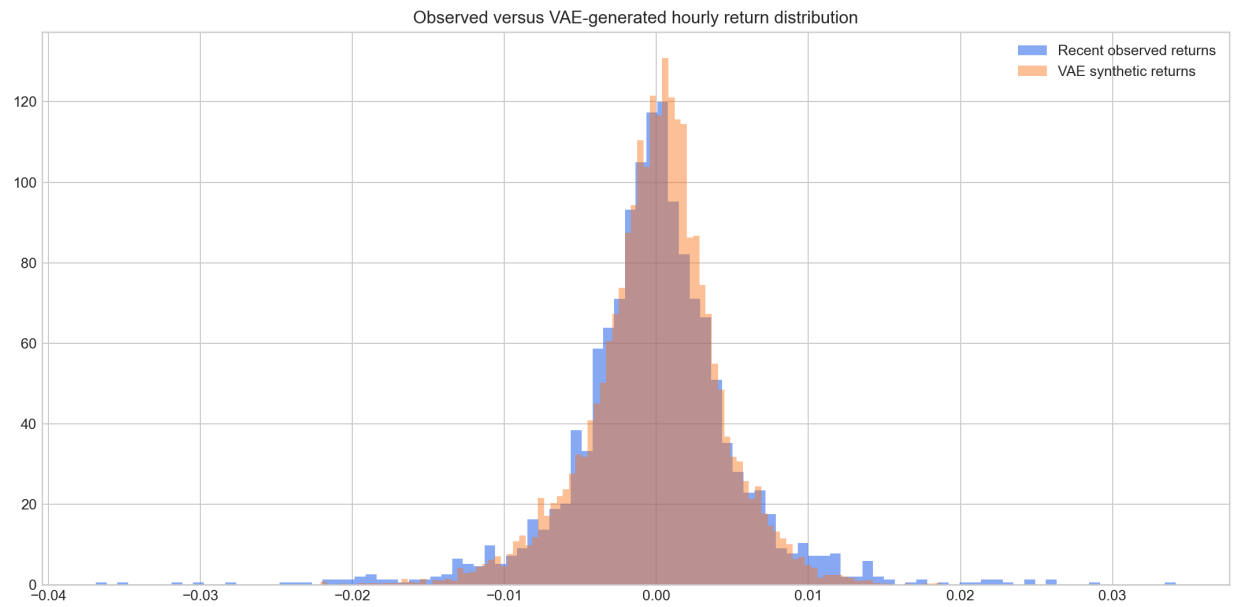

Figure 8: VAE generated returns

than the parametric baseline. Also, the validation approach is based on time rather than full rolling-origin cross-validation across several non-overlapping regimes. Finally, the VaR and CVaR data that were presented are market-risk statistics, not losses from liquidation that have been adjusted for liquidity.

Discussion

Two empirical themes predominate the evidence. Bitcoin has persistent volatility clustering, which means that state dependency is important even when raw returns are only weakly autocorrelated (Engle, 1982). Second, the downside tail is too high for a Gaussian benchmark to be believable. The Gaussian row shows less tail risk than the matching Student-t benchmark and is weaker than Student-t VaR on the final-year backtest. Those facts show why the project is better off with layers of models that give highly varied risk estimations instead of picking one winner for each use case.

Conclusion

The primary conclusion is that Gaussian hourly VaR is an insufficient baseline for Bitcoin. Compared to a thin-tailed variance-only baseline, heavy-tailed parametric modeling, jump-aware simulation, and conditional neural forecasting all give very different risk estimations. The jump-diffusion engine is the best conservative scenario benchmark for this project. The LSTM is the sole model that passed the conditional-coverage backtest for the last holdout year.

References

1. Jorion, Philippe. 2007. Value at Risk: The New Benchmark for Managing Financial Risk. 3rd ed. New York: McGraw-Hill.

2. McNeil, Alexander J., Rudiger Frey, and Paul Embrechts. 2015. Quantitative Risk Management: Concepts, Techniques and Tools. Revised ed. Princeton: Princeton University

Press.

3. Kupiec, Paul H. 1995. Techniques for Verifying the Accuracy of Risk Measurement Models. Journal of Derivatives 3(2): 73-84. DOI: 10.3905/jod.1995.407942.

4. Christoffersen, Peter F. 1998. Evaluating Interval Forecasts. International Economic Review 39(4): 841-862. DOI: 10.2307/2527341.

5. Merton, Robert C. 1976. Option Pricing When Underlying Stock Returns Are Discontinuous. Journal of Financial Economics 3(1-2): 125-144. DOI: 10.1016/0304-405X(76)90022-2.

6. Engle, Robert F. 1982. Autoregressive Conditional Heteroscedasticity with Estimates of the Variance of United Kingdom Inflation. Econometrica 50(4): 987-1007. DOI: 10.2307/1912773.

7. Hochreiter, Sepp, and Jurgen Schmidhuber. 1997. Long Short-Term Memory. Neural Computation 9(8): 1735-1780. DOI: 10.1162/neco.1997.9.8.1735.

8. Kingma, Diederik P., and Max Welling. 2014. Auto-Encoding Variational Bayes. 2nd International Conference on Learning Representations. DOI: 10.48550/arXiv.1312.6114.